



# Waffle Charts, Word Clouds, and Regression Plots

Estimated time needed: **30** minutes

## Objectives

After completing this lab you will be able to:

- Create Word cloud and Waffle charts
- Create regression plots with Seaborn library

## Table of Contents

1. [Exploring Datasets with \_p\_andas](#0)
2. [Downloading and Prepping Data](#2)
3. [Visualizing Data using Matplotlib](#4)
4. [Waffle Charts](#6)
5. [Word Clouds](#8)
6. [Regression Plots](#10)

# Exploring Datasets with *pandas* and Matplotlib

Toolkits: The course heavily relies on [pandas](http://pandas.pydata.org?cm_mmc=Email_Newsletter- -Developer_Ed%2BTech- -WW_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-) ([http://pandas.pydata.org?cm\\_mmc=Email\\_Newsletter- -Developer\\_Ed%2BTech- -WW\\_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-](http://pandas.pydata.org?cm_mmc=Email_Newsletter- -Developer_Ed%2BTech- -WW_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-))  
[Numpy](http://www.numpy.org?cm_mmc=Email_Newsletter- -Developer_Ed%2BTech- -WW_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-) ([http://www.numpy.org?cm\\_mmc=Email\\_Newsletter- -Developer\\_Ed%2BTech- -WW\\_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-](http://www.numpy.org?cm_mmc=Email_Newsletter- -Developer_Ed%2BTech- -WW_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-))  
[Matplotlib](http://matplotlib.org?cm_mmc=Email_Newsletter- -Developer_Ed%2BTech- -WW_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-) ([http://matplotlib.org?cm\\_mmc=Email\\_Newsletter- -Developer\\_Ed%2BTech- -WW\\_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-](http://matplotlib.org?cm_mmc=Email_Newsletter- -Developer_Ed%2BTech- -WW_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-))  
for data wrangling, analysis, and visualization. The primary plotting library we will explore in the course is [Matplotlib](http://www.un.org/en/development/desa/population/migration/data/empirical2/migrationflows.shtml?cm_mmc=Email_Newsletter- -Developer_Ed%2BTech- -WW_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-) ([http://www.un.org/en/development/desa/population/migration/data/empirical2/migrationflows.shtml?cm\\_mmc=Email\\_Newsletter- -Developer\\_Ed%2BTech- -WW\\_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-](http://www.un.org/en/development/desa/population/migration/data/empirical2/migrationflows.shtml?cm_mmc=Email_Newsletter- -Developer_Ed%2BTech- -WW_WW- -SkillsNetwork-Courses-IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork-)) from United Nation's website

The dataset contains annual data on the flows of international migrants as recorded by the countries of destination. The data presents both inflows and outflows according to the place of birth, citizenship or place of previous / next residence both for foreigners and nationals. In this lab, we will focus on the Canadian Immigration data.



## Downloading and Prepping Data

Import Primary Modules:

In [3]:

```
import numpy as np # useful for many scientific computing in Python
import pandas as pd # primary data structure library
from PIL import Image # converting images into arrays
```

Let's download and import our primary Canadian Immigration dataset using *pandas read\_excel()* method. Normally, before we can do that, we would need to download a module which *pandas* requires to read in excel files. This module is **xlrd**. For your convenience, we have pre-installed this module, so you would not have to worry about that. Otherwise, you would need to run the following line of code to install the **xlrd** module:

```
!conda install -c anaconda xlrd --yes
```

Download the dataset and read it into a *pandas* dataframe:

In [5]:

```
df_can = pd.read_excel('https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDriverSkillsNetwork-DV0101EN-SkillsNetwork/Data%20Files/Canada.xlsx',
                       sheet_name='Canada by Citizenship',
                       skiprows=range(20),
                       skipfooter=2)

print('Data downloaded and read into a dataframe!')
```

Data downloaded and read into a dataframe!

Let's take a look at the first five items in our dataset

In [6]:

```
df_can.head()
```

Out[6]:

	Type	Coverage	OdName	AREA	AreaName	REG	RegName	DEV	DevName	1980
0	Immigrants	Foreigners	Afghanistan	935	Asia	5501	Southern Asia	902	Developing regions	16
1	Immigrants	Foreigners	Albania	908	Europe	925	Southern Europe	901	Developed regions	1
2	Immigrants	Foreigners	Algeria	903	Africa	912	Northern Africa	902	Developing regions	80
3	Immigrants	Foreigners	American Samoa	909	Oceania	957	Polynesia	902	Developing regions	0
4	Immigrants	Foreigners	Andorra	908	Europe	925	Southern Europe	901	Developed regions	0

5 rows × 43 columns



Let's find out how many entries there are in our dataset

In [7]:

```
# print the dimensions of the dataframe
print(df_can.shape)
```

(195, 43)

Clean up data. We will make some modifications to the original dataset to make it easier to create our visualizations. Refer to *Introduction to Matplotlib and Line Plots and Area Plots, Histograms, and Bar Plots* for a detailed description of this preprocessing.

In [8]:

```
# clean up the dataset to remove unnecessary columns (eg. REG)
df_can.drop(['AREA', 'REG', 'DEV', 'Type', 'Coverage'], axis = 1, inplace = True)

# Let's rename the columns so that they make sense
df_can.rename (columns = {'OdName':'Country', 'AreaName':'Continent', 'RegName':'Region'}, inplace = True)

# for sake of consistency, let's also make all column labels of type string
df_can.columns = list(map(str, df_can.columns))

# set the country name as index - useful for quickly looking up countries using .loc method
df_can.set_index('Country', inplace = True)

# add total column
df_can['Total'] = df_can.sum (axis = 1)

# years that we will be using in this lesson - useful for plotting later on
years = list(map(str, range(1980, 2014)))
print ('data dimensions:', df_can.shape)
```

data dimensions: (195, 38)

## Visualizing Data using Matplotlib

Import matplotlib :

In [9]:

```
%matplotlib inline

import matplotlib as mpl
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches # needed for waffle Charts

mpl.style.use('ggplot') # optional: for ggplot-like style

# check for latest version of Matplotlib
print ('Matplotlib version: ', mpl.__version__) # >= 2.0.0
```

Matplotlib version: 3.3.3

## Waffle Charts

A waffle chart is an interesting visualization that is normally created to display progress toward goals. It is commonly an effective option when you are trying to add interesting visualization features to a visual that consists mainly of cells, such as an Excel dashboard.

Let's revisit the previous case study about Denmark, Norway, and Sweden.

In [10]:

```
# Let's create a new dataframe for these three countries
df_dsn = df_can.loc[['Denmark', 'Norway', 'Sweden'], :]

# Let's take a look at our dataframe
df_dsn
```

Out[10]:

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	:
Country													
<b>Denmark</b>	Europe	Northern Europe	Developed regions	272	293	299	106	93	73	93	...	62	
<b>Norway</b>	Europe	Northern Europe	Developed regions	116	77	106	51	31	54	56	...	57	
<b>Sweden</b>	Europe	Northern Europe	Developed regions	281	308	222	176	128	158	187	...	205	

3 rows × 38 columns



Unfortunately, unlike R, waffle charts are not built into any of the Python visualization libraries. Therefore, we will learn how to create them from scratch.

**Step 1.** The first step into creating a waffle chart is determing the proportion of each category with respect to the total.

In [11]:

```
# compute the proportion of each category with respect to the total
total_values = sum(df_dsn['Total'])
category_proportions = [(float(value) / total_values) for value in df_dsn['Total']]

# print out proportions
for i, proportion in enumerate(category_proportions):
    print (df_dsn.index.values[i] + ': ' + str(proportion))
```

Denmark: 0.32255663965602777  
Norway: 0.1924094592359848  
Sweden: 0.48503390110798744

**Step 2.** The second step is defining the overall size of the waffle chart.

In [15]:

```
width = 40 # width of chart
height = 10 # height of chart

total_num_tiles = width * height # total number of tiles

print ('Total number of tiles is ', total_num_tiles)
```

Total number of tiles is 400

**Step 3.** The third step is using the proportion of each category to determe it respective number of tiles

In [17]:

```
# compute the number of tiles for each catagory
tiles_per_category = [round(proportion * total_num_tiles) for proportion in category_proportions]

# print out number of tiles per category
for i, tiles in enumerate(tiles_per_category):
    print (df_dsn.index.values[i] + ': ' + str(tiles))
```

Denmark: 129  
Norway: 77  
Sweden: 194

Based on the calculated proportions, Denmark will occupy 129 tiles of the waffle chart, Norway will occupy 77 tiles, and Sweden will occupy 194 tiles.

**Step 4.** The fourth step is creating a matrix that resembles the waffle chart and populating it.

In [18]:

```
# initialize the waffle chart as an empty matrix
waffle_chart = np.zeros((height, width))

# define indices to loop through waffle chart
category_index = 0
tile_index = 0

# populate the waffle chart
for col in range(width):
    for row in range(height):
        tile_index += 1

        # if the number of tiles populated for the current category is equal to its corresponding allocated tiles...
        if tile_index > sum(tiles_per_category[0:category_index]):
            # ...proceed to the next category
            category_index += 1

        # set the class value to an integer, which increases with class
        waffle_chart[row, col] = category_index

print ('Waffle chart populated!')
```

Waffle chart populated!

Let's take a peek at how the matrix looks like.

In [19]:

## waffle\_chart

Out[19]:

As expected, the matrix consists of three categories and the total number of each category's instances matches the total number of tiles allocated to each category.

**Step 5.** Map the waffle chart matrix into a visual.

In [20]:

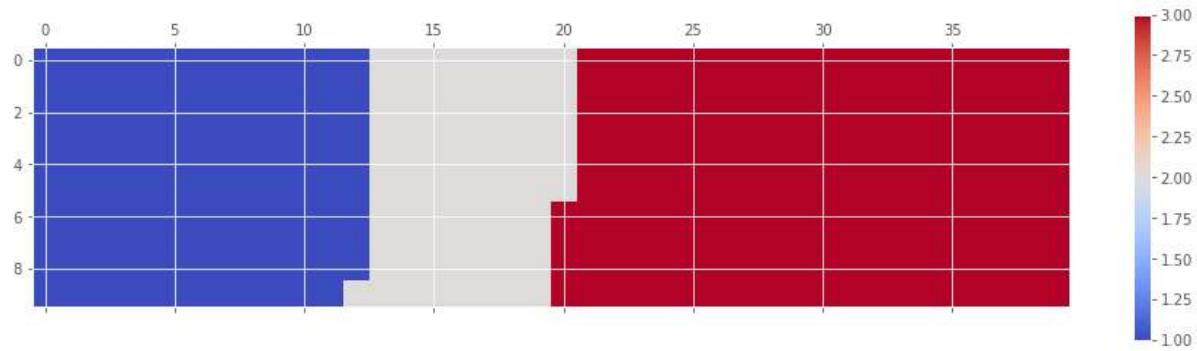
```
# instantiate a new figure object
fig = plt.figure()

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()
```

Out[20]:

```
<matplotlib.colorbar.Colorbar at 0x7fb3dc6894a8>
```

```
<Figure size 432x288 with 0 Axes>
```



**Step 6.** Prettify the chart.

In [21]:

```
# instantiate a new figure object
fig = plt.figure()

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()

# get the axis
ax = plt.gca()

# set minor ticks
ax.set_xticks(np.arange(-.5, (width), 1), minor=True)
ax.set_yticks(np.arange(-.5, (height), 1), minor=True)

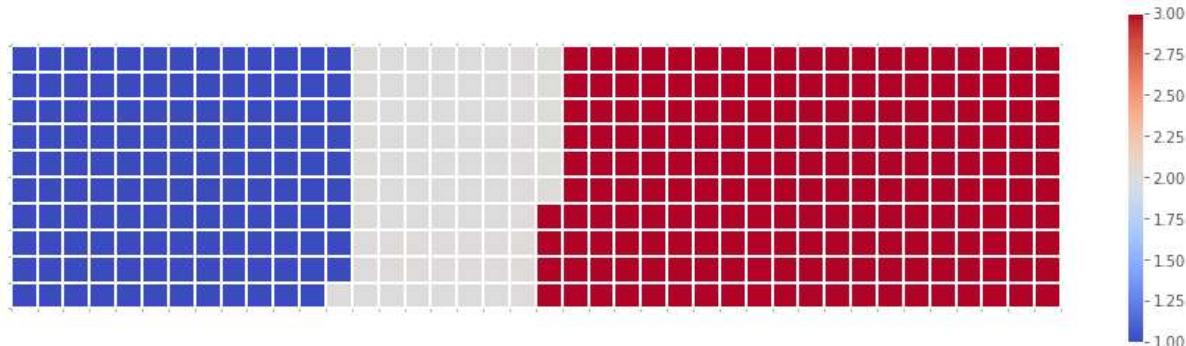
# add gridlines based on minor ticks
ax.grid(which='minor', color='w', linestyle='-', linewidth=2)

plt.xticks([])
plt.yticks([])
```

Out[21]:

([], [])

<Figure size 432x288 with 0 Axes>



**Step 7.** Create a legend and add it to chart.

In [22]:

```
# instantiate a new figure object
fig = plt.figure()

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()

# get the axis
ax = plt.gca()

# set minor ticks
ax.set_xticks(np.arange(-.5, (width), 1), minor=True)
ax.set_yticks(np.arange(-.5, (height), 1), minor=True)

# add gridlines based on minor ticks
ax.grid(which='minor', color='w', linestyle='-', linewidth=2)

plt.xticks([])
plt.yticks([])

# compute cumulative sum of individual categories to match color schemes between chart and
# legend
values_cumsum = np.cumsum(df_dsn['Total'])
total_values = values_cumsum[len(values_cumsum) - 1]

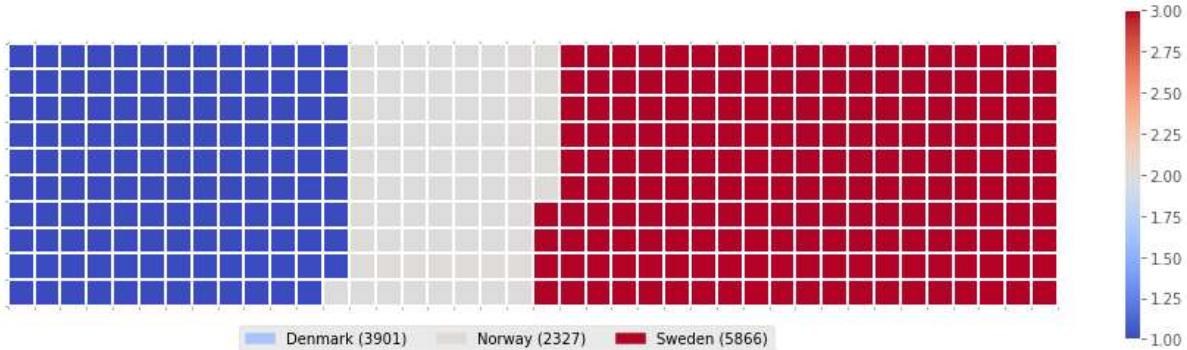
# create legend
legend_handles = []
for i, category in enumerate(df_dsn.index.values):
    label_str = category + ' (' + str(df_dsn['Total'][i]) + ')'
    color_val = colormap(float(values_cumsum[i])/total_values)
    legend_handles.append(mpatches.Patch(color=color_val, label=label_str))

# add legend to chart
plt.legend(handles=legend_handles,
           loc='lower center',
           ncol=len(df_dsn.index.values),
           bbox_to_anchor=(0., -0.2, 0.95, .1)
)
```

Out[22]:

```
<matplotlib.legend.Legend at 0x7fb3dc34cb70>
```

```
<Figure size 432x288 with 0 Axes>
```



And there you go! What a good looking *delicious* waffle chart, don't you think?

Now it would very inefficient to repeat these seven steps every time we wish to create a `waffle` chart. So let's combine all seven steps into one function called `_create_wafflechart`. This function would take the following parameters as input:

1. **categories**: Unique categories or classes in dataframe.
2. **values**: Values corresponding to categories or classes.
3. **height**: Defined height of waffle chart.
4. **width**: Defined width of waffle chart.
5. **colormap**: Colormap class
6. **value\_sign**: In order to make our function more generalizable, we will add this parameter to address signs that could be associated with a value such as %, \$, and so on. **value\_sign** has a default value of empty string.

In [23]:

```
def create_waffle_chart(categories, values, height, width, colormap, value_sign=''):

    # compute the proportion of each category with respect to the total
    total_values = sum(values)
    category_proportions = [(float(value) / total_values) for value in values]

    # compute the total number of tiles
    total_num_tiles = width * height # total number of tiles
    print ('Total number of tiles is', total_num_tiles)

    # compute the number of tiles for each category
    tiles_per_category = [round(proportion * total_num_tiles) for proportion in category_proportions]

    # print out number of tiles per category
    for i, tiles in enumerate(tiles_per_category):
        print (df_dsn.index.values[i] + ': ' + str(tiles))

    # initialize the waffle chart as an empty matrix
    waffle_chart = np.zeros((height, width))

    # define indices to loop through waffle chart
    category_index = 0
    tile_index = 0

    # populate the waffle chart
    for col in range(width):
        for row in range(height):
            tile_index += 1

            # if the number of tiles populated for the current category
            # is equal to its corresponding allocated tiles...
            if tile_index > sum(tiles_per_category[0:category_index]):
                # ...proceed to the next category
                category_index += 1

            # set the class value to an integer, which increases with class
            waffle_chart[row, col] = category_index

    # instantiate a new figure object
    fig = plt.figure()

    # use matshow to display the waffle chart
    colormap = plt.cm.coolwarm
    plt.matshow(waffle_chart, cmap=colormap)
    plt.colorbar()

    # get the axis
    ax = plt.gca()

    # set minor ticks
    ax.set_xticks(np.arange(-.5, (width), 1), minor=True)
    ax.set_yticks(np.arange(-.5, (height), 1), minor=True)

    # add gridlines based on minor ticks
```

```

ax.grid(which='minor', color='w', linestyle='-', linewidth=2)

plt.xticks([])
plt.yticks([])

# compute cumulative sum of individual categories to match color schemes between chart and legend
values_cumsum = np.cumsum(values)
total_values = values_cumsum[len(values_cumsum) - 1]

# create legend
legend_handles = []
for i, category in enumerate(categories):
    if value_sign == '%':
        label_str = category + ' (' + str(values[i]) + value_sign + ')'
    else:
        label_str = category + ' (' + value_sign + str(values[i]) + ')'

    color_val = colormap(float(values_cumsum[i])/total_values)
    legend_handles.append(mpatches.Patch(color=color_val, label=label_str))

# add Legend to chart
plt.legend(
    handles=legend_handles,
    loc='lower center',
    ncol=len(categories),
    bbox_to_anchor=(0., -0.2, 0.95, .1)
)

```

Now to create a `waffle` chart, all we have to do is call the function `create_waffle_chart`. Let's define the input parameters:

In [24]:

```

width = 40 # width of chart
height = 10 # height of chart

categories = df_dsn.index.values # categories
values = df_dsn['Total'] # correponding values of categories

colormap = plt.cm.coolwarm # color map class

```

And now let's call our function to create a `waffle` chart.

In [25]:

```
create_waffle_chart(categories, values, height, width, colormap)
```

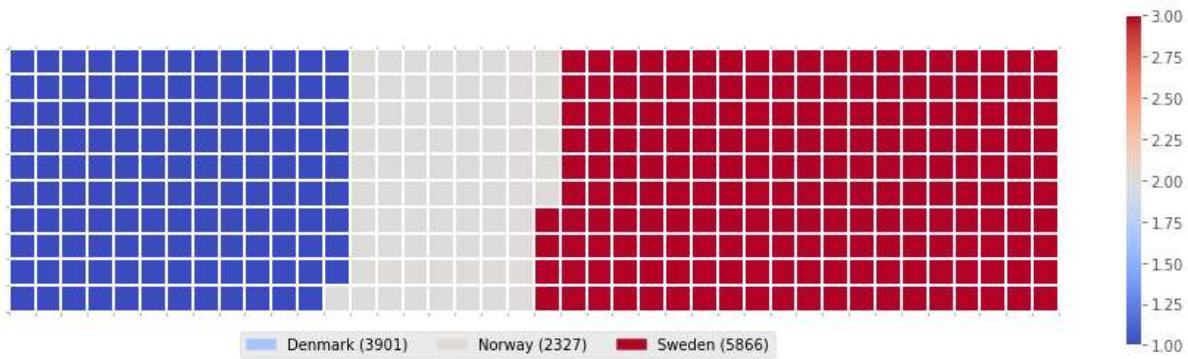
Total number of tiles is 400

Denmark: 129

Norway: 77

Sweden: 194

<Figure size 432x288 with 0 Axes>



There seems to be a new Python package for generating waffle charts called [PyWaffle](#) (<https://github.com/ligyxy/PyWaffle>), but it looks like the repository is still being built. But feel free to check it out and play with it.

## Word Clouds

Word clouds (also known as text clouds or tag clouds) work in a simple way: the more a specific word appears in a source of textual data (such as a speech, blog post, or database), the bigger and bolder it appears in the word cloud.

Luckily, a Python package already exists in Python for generating word clouds. The package, called `word_cloud` was developed by **Andreas Mueller**. You can learn more about the package by following this [link](#) ([https://github.com/amueller/word\\_cloud/](https://github.com/amueller/word_cloud/)).

Let's use this package to learn how to generate a word cloud for a given text document.

First, let's install the package.

In [28]:

```
# install wordcloud
!conda install -c conda-forge wordcloud==1.4.1 --yes

# import package and its set of stopwords
from wordcloud import WordCloud, STOPWORDS

print ('Wordcloud is installed and imported!')
```

Collecting package metadata (current\_repodata.json): done  
Solving environment: done

# All requested packages already installed.

Wordcloud is installed and imported!

Word clouds are commonly used to perform high-level analysis and visualization of text data. Accordingly, let's digress from the immigration dataset and work with an example that involves analyzing text data. Let's try to analyze a short novel written by **Lewis Carroll** titled *Alice's Adventures in Wonderland*. Let's go ahead and download a .txt file of the novel.

In [29]:

```
# download file and save as alice_novel.txt
!wget --quiet https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DV0101EN-SkillsNetwork/Data%20Files/alice_novel.txt

# open the file and read it into a variable alice_novel
alice_novel = open('alice_novel.txt', 'r').read()

print ('File downloaded and saved!')
```

File downloaded and saved!

Next, let's use the stopwords that we imported from `word_cloud`. We use the function `set` to remove any redundant stopwords.

In [30]:

```
stopwords = set(STOPWORDS)
```

Create a word cloud object and generate a word cloud. For simplicity, let's generate a word cloud using only the first 2000 words in the novel.

In [31]:

```
# instantiate a word cloud object
alice_wc = WordCloud(
    background_color='white',
    max_words=2000,
    stopwords=stopwords
)

# generate the word cloud
alice_wc.generate(alice_novel)
```

Out[31]:

```
<wordcloud.wordcloud.WordCloud at 0x7fb33eb61550>
```

Awesome! Now that the word cloud is created, let's visualize it.

In [32]:

```
# display the word cloud
plt.imshow(alice_wc, interpolation='bilinear')
plt.axis('off')
plt.show()
```



Interesting! So in the first 2000 words in the novel, the most common words are **Alice**, **said**, **little**, **Queen**, and so on. Let's resize the cloud so that we can see the less frequent words a little better.

In [33]:

```
fig = plt.figure()
fig.set_figwidth(14) # set width
fig.set_figheight(18) # set height

# display the cloud
plt.imshow(alice_wc, interpolation='bilinear')
plt.axis('off')
plt.show()
```



Much better! However, **said** isn't really an informative word. So let's add it to our stopwords and re-generate the cloud.

In [34]:

```
stopwords.add('said') # add the words said to stopwords

# re-generate the word cloud
alice_wc.generate(alice_novel)

# display the cloud
fig = plt.figure()
fig.set_figwidth(14) # set width
fig.set_figheight(18) # set height

plt.imshow(alice_wc, interpolation='bilinear')
plt.axis('off')
plt.show()
```



Excellent! This looks really interesting! Another cool thing you can implement with the `word_cLOUD` package is superimposing the words onto a mask of any shape. Let's use a mask of Alice and her rabbit. We already created the mask for you, so let's go ahead and download it and call it `_alicemask.png`.

In [35]:

```
# download image
!wget --quiet https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DV0101EN-SkillsNetwork/labs/Module%204/images/alice_mask.png

# save mask to alice_mask
alice_mask = np.array(Image.open('alice_mask.png'))

print('Image downloaded and saved!')
```

Image downloaded and saved!

Let's take a look at how the mask looks like.

In [36]:

```
fig = plt.figure()
fig.set_figwidth(14) # set width
fig.set_figheight(18) # set height

plt.imshow(alice_mask, cmap=plt.cm.gray, interpolation='bilinear')
plt.axis('off')
plt.show()
```



Shaping the word cloud according to the mask is straightforward using `word_cloud` package. For simplicity, we will continue using the first 2000 words in the novel.

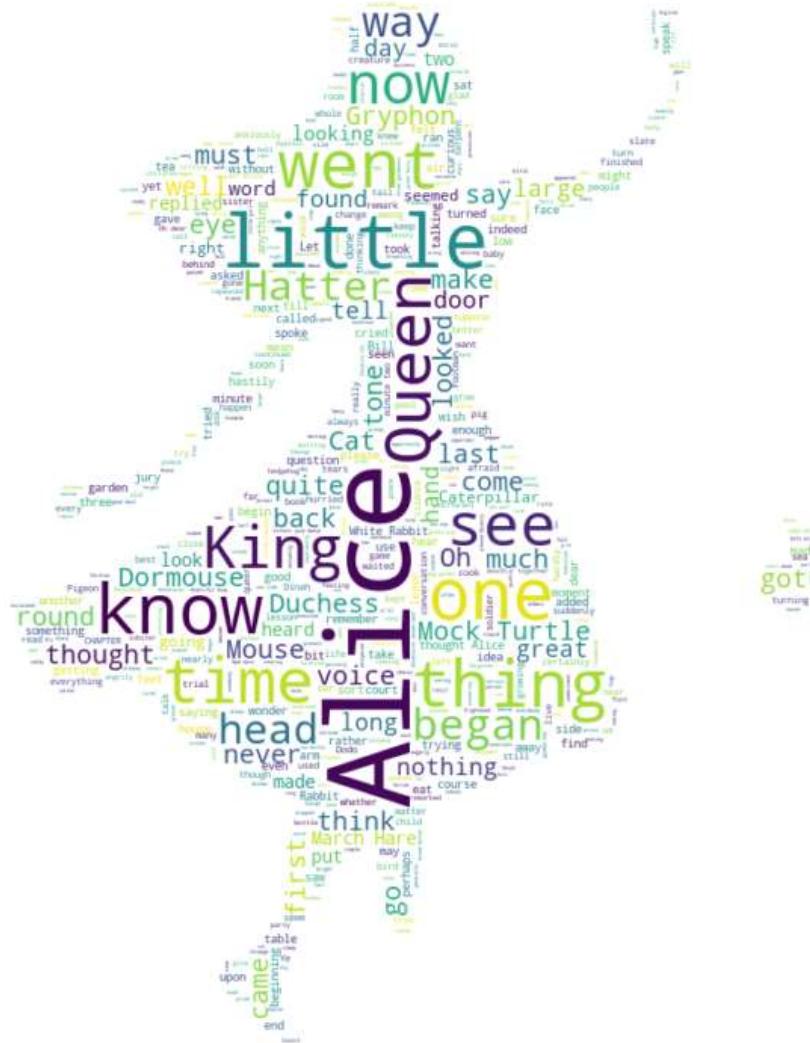
In [37]:

```
# instantiate a word cloud object
alice_wc = WordCloud(background_color='white', max_words=2000, mask=alice_mask, stopwords=
stopwords)

# generate the word cloud
alice_wc.generate(alice_novel)

# display the word cloud
fig = plt.figure()
fig.set_figwidth(14) # set width
fig.set_figheight(18) # set height

plt.imshow(alice_wc, interpolation='bilinear')
plt.axis('off')
plt.show()
```



Really impressive!

Unfortunately, our immigration data does not have any text data, but where there is a will there is a way. Let's generate sample text data from our immigration dataset, say text data of 90 words.

Let's recall how our data looks like.

In [38]:

```
df_can.head()
```

Out[38]:

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2013
Country												
Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	496	...	341
Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	1	...	12
Algeria	Africa	Northern Africa	Developing regions	80	67	71	69	63	44	69	...	36
American Samoa	Oceania	Polynesia	Developing regions	0	1	0	0	0	0	0	...	
Andorra	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	2	...	

5 rows × 38 columns



And what was the total immigration from 1980 to 2013?

In [39]:

```
total_immigration = df_can['Total'].sum()
total_immigration
```

Out[39]:

6409153

Using countries with single-word names, let's duplicate each country's name based on how much they contribute to the total immigration.

In [40]:

```
max_words = 90
word_string = ''
for country in df_can.index.values:
    # check if country's name is a single-word name
    if len(country.split(' ')) == 1:
        repeat_num_times = int(df_can.loc[country, 'Total'])/float(total_immigration)*max_words
        word_string = word_string + ((country + ' ') * repeat_num_times)

# display the generated text
word_string
```

Out[40]:

```
'China China China China China China China Colombia Egypt France
Guyana Haiti India India India India India India India Jamaica Le
banon Morocco Pakistan Pakistan Philippines Philippines Philippines
Philippines Philippines Philippines Poland Portugal Romania '
```

We are not dealing with any stopwords here, so there is no need to pass them when creating the word cloud.

In [41]:

```
# create the word cloud
wordcloud = WordCloud(background_color='white').generate(word_string)

print('Word cloud created!')
```

Word cloud created!

In [42]:

```
# display the cloud
fig = plt.figure()
fig.set_figwidth(14)
fig.set_figheight(18)

plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.show()
```



According to the above word cloud, it looks like the majority of the people who immigrated came from one of 15 countries that are displayed by the word cloud. One cool visual that you could build, is perhaps using the map of Canada and a mask and superimposing the word cloud on top of the map of Canada. That would be an interesting visual to build!

# Regression Plots



In lab *Pie Charts, Box Plots, Scatter Plots, and Bubble Plots*, we learned how to create a scatter plot and then fit a regression line. It took ~20 lines of code to create the scatter plot along with the regression fit. In this final section, we will explore *seaborn* and see how efficient it is to create regression lines and fits using this library!

Let's first install *seaborn*

In [43]:

```
# install seaborn
!conda install -c anaconda seaborn --yes

# import library
import seaborn as sns

print('Seaborn installed and imported!')
```

```
Collecting package metadata (current_repodata.json): done
Solving environment: done
```

```
## Package Plan ##
```

```
environment location: /home/jupyterlab/conda/envs/python
```

```
added / updated specs:
```

```
- seaborn
```

The following packages will be downloaded:

package	build			
blas-1.0	mkl	6 KB	anaconda	
ca-certificates-2020.10.14	0	128 KB	anaconda	
certifi-2020.6.20	py36_0	160 KB	anaconda	
dbus-1.13.16	hb2f20db_0	589 KB	anaconda	
gst-plugins-base-1.14.0	hbbd80ab_1	4.8 MB		
gstreamer-1.14.0	h28cd5cc_2	3.2 MB		
libgfortran-ng-7.3.0	hdf63c60_0	1.3 MB	anaconda	
matplotlib-3.3.1	0	24 KB	anaconda	
matplotlib-base-3.3.1	py36h817c723_0	6.7 MB	anaconda	
mkl-service-2.3.0	py36he8ac12f_0	52 KB		
mkl_fft-1.2.0	py36h23d657b_0	149 KB		
mkl_random-1.1.1	py36h0573a6f_0	382 KB	anaconda	
numpy-1.19.2	py36h54aff64_0	22 KB		
numpy-base-1.19.2	py36hfa32c7d_0	4.1 MB		
openssl-1.1.1h	h7b6447c_0	3.8 MB	anaconda	
pandas-1.1.3	py36he6710b0_0	10.5 MB	anaconda	
pyqt-5.9.2	py36h22d08a2_1	5.6 MB	anaconda	
pytz-2020.1	py_0	239 KB	anaconda	
qt-5.9.7	h5867ecd_1	68.5 MB		
scipy-1.5.2	py36h0b6359f_0	14.4 MB		
seaborn-0.11.0	py_0	216 KB	anaconda	
sip-4.19.24	py36he6710b0_0	297 KB	anaconda	
<hr/>				
			Total:	125.2 MB

The following NEW packages will be INSTALLED:

blas	anaconda/linux-64::blas-1.0-mkl
dbus	anaconda/linux-64::dbus-1.13.16-hb2f20db_0
gst-plugins-base	pkgs/main/linux-64::gst-plugins-base-1.14.0-hbbd80ab_1
gstreamer	pkgs/main/linux-64::gstreamer-1.14.0-h28cd5cc_2
matplotlib	anaconda/linux-64::matplotlib-3.3.1-0
mkl-service	pkgs/main/linux-64::mkl-service-2.3.0-py36he8ac12f_0
mkl_fft	pkgs/main/linux-64::mkl_fft-1.2.0-py36h23d657b_0
mkl_random	anaconda/linux-64::mkl_random-1.1.1-py36h0573a6f_0
numpy-base	pkgs/main/linux-64::numpy-base-1.19.2-py36hfa32c7d_0
pandas	anaconda/linux-64::pandas-1.1.3-py36he6710b0_0
pyqt	anaconda/linux-64::pyqt-5.9.2-py36h22d08a2_1
pytz	anaconda/noarch::pytz-2020.1-py_0
qt	pkgs/main/linux-64::qt-5.9.7-h5867ecd_1
scipy	pkgs/main/linux-64::scipy-1.5.2-py36h0b6359f_0
seaborn	anaconda/noarch::seaborn-0.11.0-py_0

```
sip anaconda/linux-64::sip-4.19.24-py36he6710b0_0
```

The following packages will be REMOVED:

```
libblas-3.9.0-3_openblas
libcblas-3.9.0-3_openblas
libgfortran5-9.3.0-he4bcb1c_17
liblapack-3.9.0-3_openblas
libopenblas-0.3.12-pthreads_h4812303_1
```

The following packages will be SUPERSEDED by a higher-priority channel:

```
ca-certificates      conda-forge::ca-certificates-2020.12.~ --> anaconda::ca-
certificates-2020.10.14-0
certifi              conda-forge::certifi-2020.12.5-py36h5~ --> anaconda::cer-
tifi-2020.6.20-py36_0
libgfortran-ng       conda-forge::libgfortran-ng-9.3.0-he4~ --> anaconda::lib-
gfortran-ng-7.3.0-hdf63c60_0
matplotlib-base     conda-forge::matplotlib-base-3.3.3-py~ --> anaconda::mat-
plotlib-base-3.3.1-py36h817c723_0
numpy                conda-forge::numpy-1.19.4-py36h2aa4a0~ --> pkgs/main::nu-
mpy-1.19.2-py36h54aff64_0
openssl              conda-forge::openssl-1.1.1i-h7f98852_0 --> anaconda::ope-
nssl-1.1.1h-h7b6447c_0
```

#### Downloading and Extracting Packages

pytz-2020.1	239 KB	#####	10
0%			
pandas-1.1.3	10.5 MB	#####	10
0%			
seaborn-0.11.0	216 KB	#####	10
0%			
scipy-1.5.2	14.4 MB	#####	10
0%			
certifi-2020.6.20	160 KB	#####	10
0%			
qt-5.9.7	68.5 MB	#####	10
0%			
pyqt-5.9.2	5.6 MB	#####	10
0%			
blas-1.0	6 KB	#####	10
0%			
mkl-service-2.3.0	52 KB	#####	10
0%			
dbus-1.13.16	589 KB	#####	10
0%			
mkl_random-1.1.1	382 KB	#####	10
0%			
mkl_fft-1.2.0	149 KB	#####	10
0%			
matplotlib-3.3.1	24 KB	#####	10
0%			
openssl-1.1.1h	3.8 MB	#####	10
0%			
matplotlib-base-3.3.	6.7 MB	#####	10

```
0%  
numpy-1.19.2 | 22 KB | ##### | 10  
0%  
libgfortran-ng-7.3.0 | 1.3 MB | ##### | 10  
0%  
sip-4.19.24 | 297 KB | ##### | 10  
0%  
gst-plugins-base-1.1 | 4.8 MB | ##### | 10  
0%  
numpy-base-1.19.2 | 4.1 MB | ##### | 10  
0%  
gstreamer-1.14.0 | 3.2 MB | ##### | 10  
0%  
ca-certificates-2020 | 128 KB | ##### | 10  
0%  
Preparing transaction: done  
Verifying transaction: done  
Executing transaction: done  
Seaborn installed and imported!
```

Create a new dataframe that stores that total number of landed immigrants to Canada per year from 1980 to 2013.

In [44]:

```
# we can use the sum() method to get the total population per year
df_tot = pd.DataFrame(df_can[years].sum(axis=0))

# change the years to type float (useful for regression later on)
df_tot.index = map(float, df_tot.index)

# reset the index to put it back in as a column in the df_tot dataframe
df_tot.reset_index(inplace=True)

# rename columns
df_tot.columns = ['year', 'total']

# view the final dataframe
df_tot.head()
```

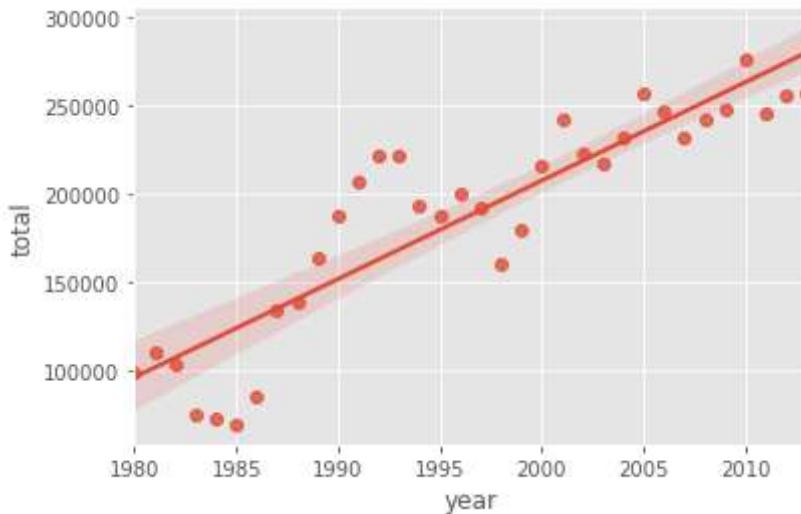
Out[44]:

	year	total
0	1980.0	99137
1	1981.0	110563
2	1982.0	104271
3	1983.0	75550
4	1984.0	73417

With `seaborn`, generating a regression plot is as simple as calling the `regplot` function.

In [45]:

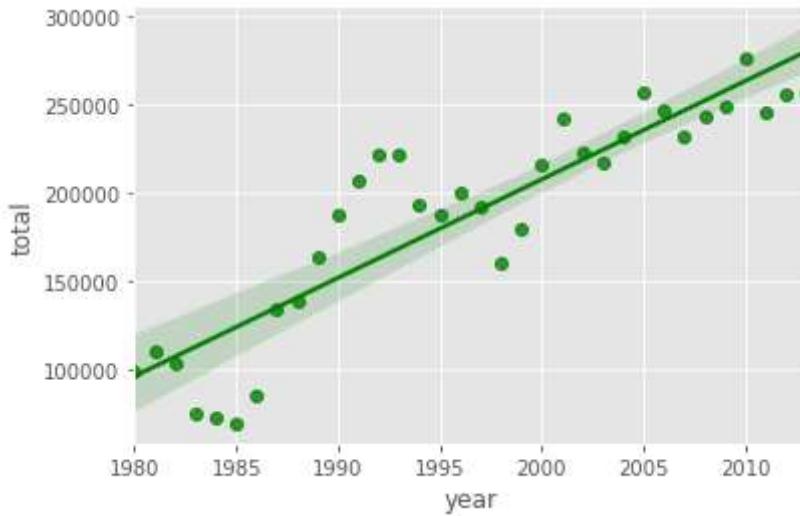
```
import seaborn as sns
ax = sns.regplot(x='year', y='total', data=df_tot)
```



This is not magic; it is `seaborn`! You can also customize the color of the scatter plot and regression line. Let's change the color to green.

In [46]:

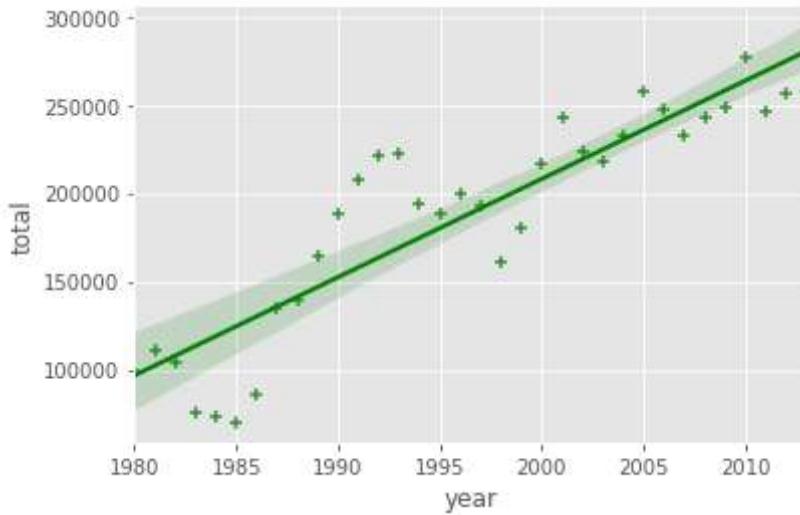
```
import seaborn as sns  
ax = sns.regplot(x='year', y='total', data=df_tot, color='green')
```



You can always customize the marker shape, so instead of circular markers, let's use '+'.

In [47]:

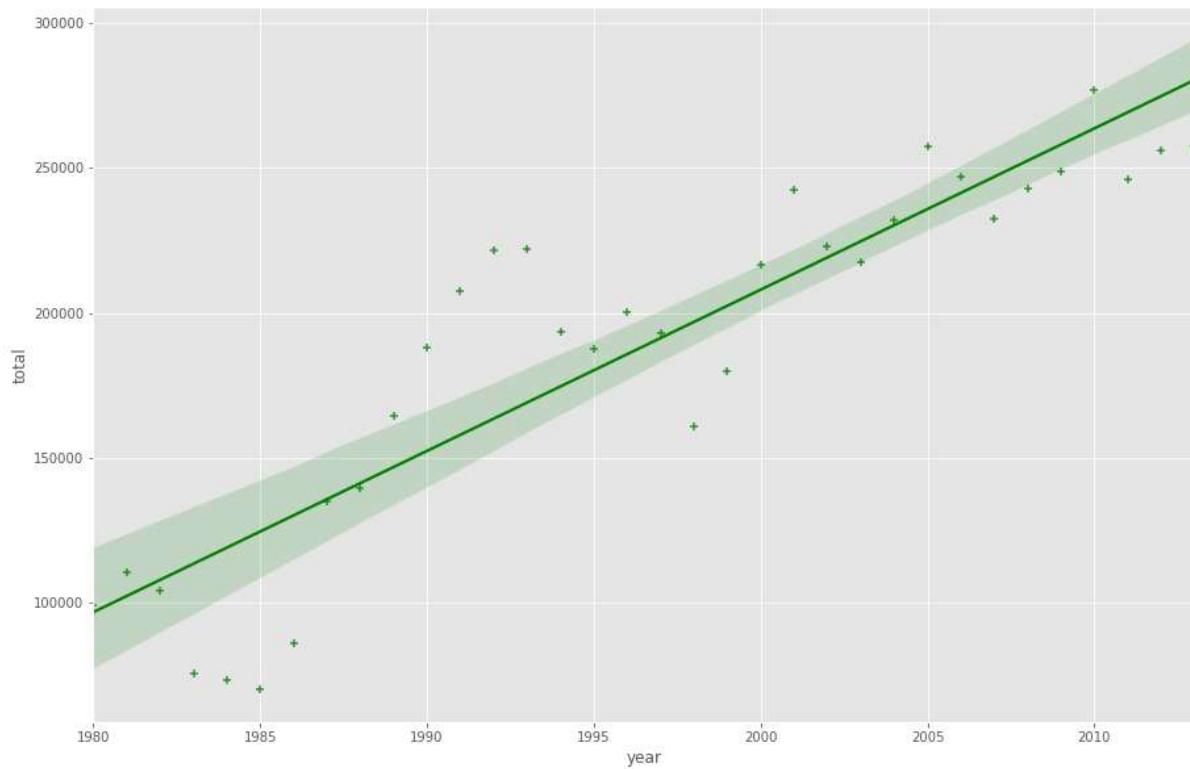
```
import seaborn as sns  
ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+')
```



Let's blow up the plot a little bit so that it is more appealing to the sight.

In [48]:

```
plt.figure(figsize=(15, 10))
ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+')
```



And let's increase the size of markers so they match the new size of the figure, and add a title and x- and y-labels.

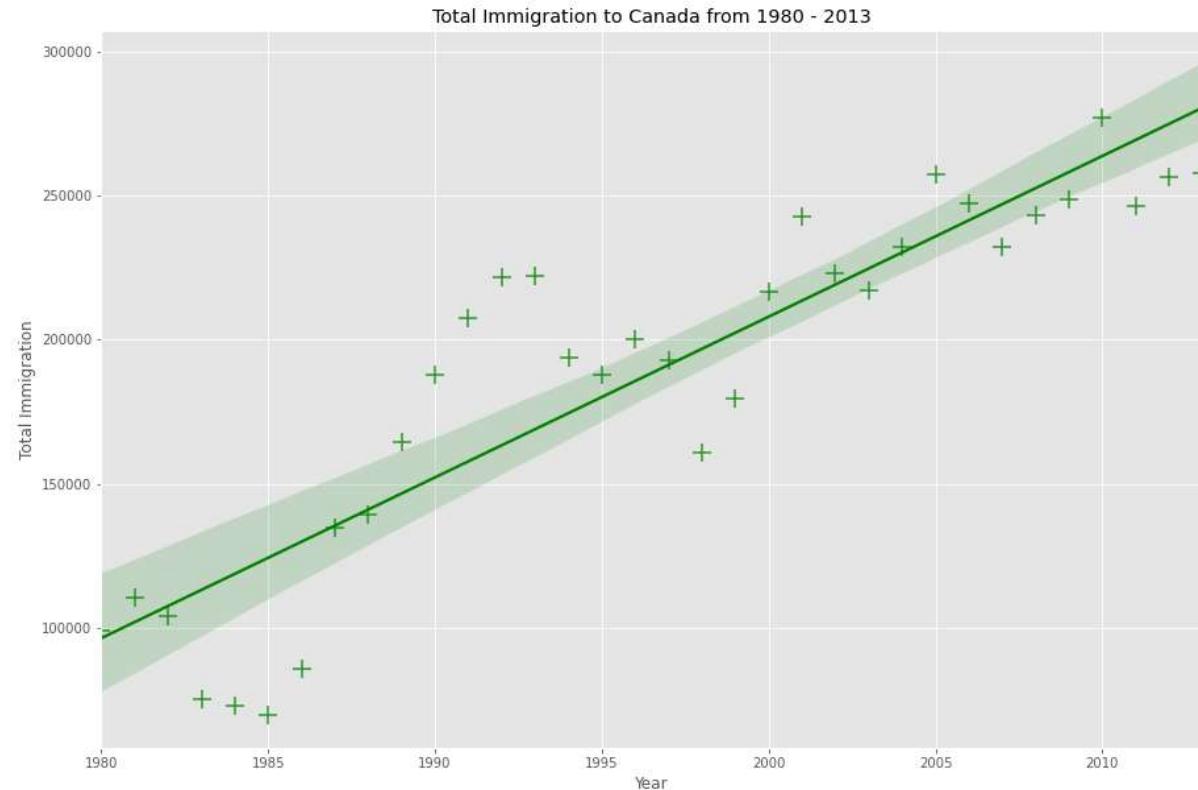
In [49]:

```
plt.figure(figsize=(15, 10))
ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+', scatter_kws=
{'s': 200})

ax.set(xlabel='Year', ylabel='Total Immigration') # add x- and y-labels
ax.set_title('Total Immigration to Canada from 1980 - 2013') # add title
```

Out[49]:

Text(0.5, 1.0, 'Total Immigration to Canada from 1980 - 2013')



And finally increase the font size of the tickmark labels, the title, and the x- and y-labels so they don't feel left out!

In [50]:

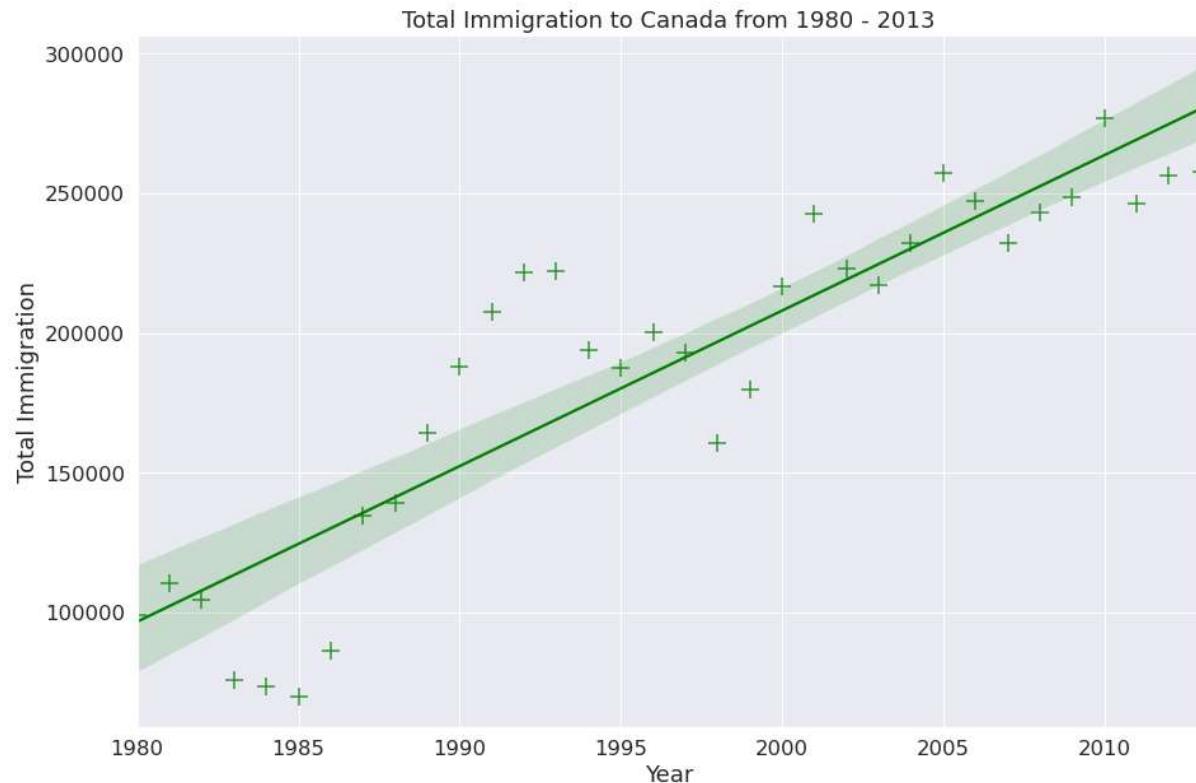
```
plt.figure(figsize=(15, 10))

sns.set(font_scale=1.5)

ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+', scatter_kws=
{'s': 200})
ax.set(xlabel='Year', ylabel='Total Immigration')
ax.set_title('Total Immigration to Canada from 1980 - 2013')
```

Out[50]:

Text(0.5, 1.0, 'Total Immigration to Canada from 1980 - 2013')



Amazing! A complete scatter plot with a regression fit with 5 lines of code only. Isn't this really amazing?

If you are not a big fan of the purple background, you can easily change the style to a white plain background.

In [51]:

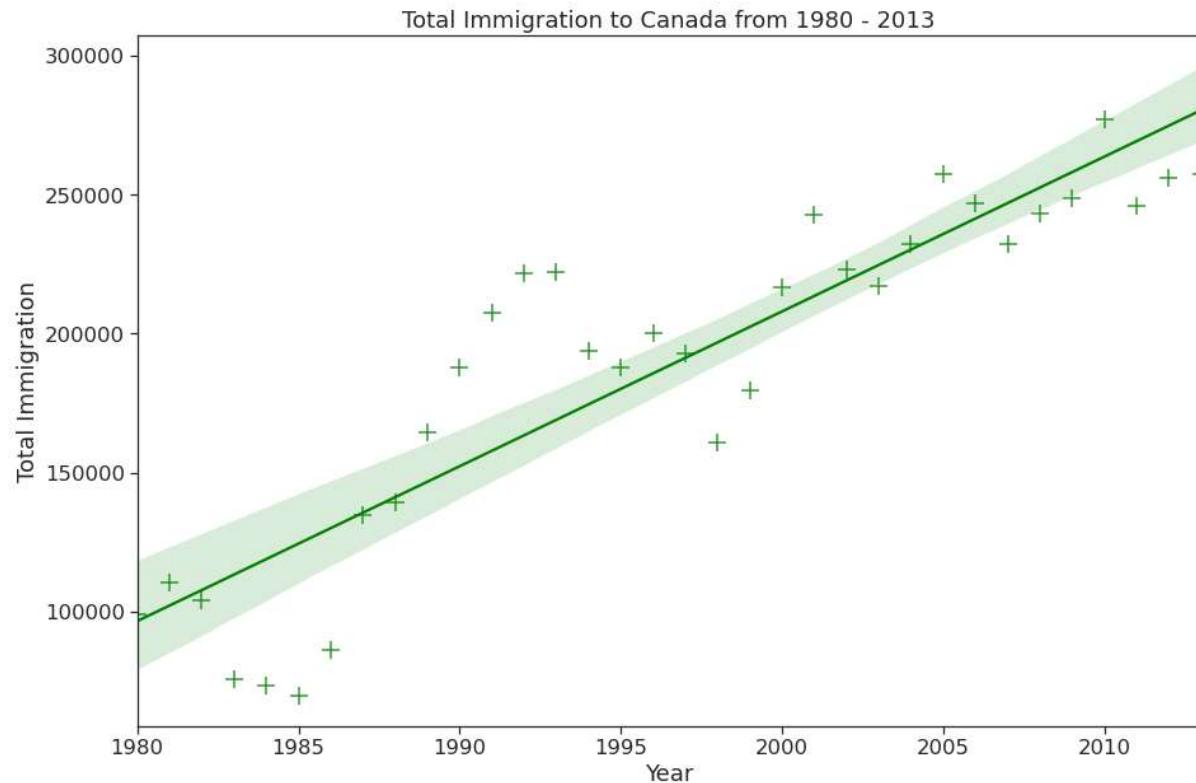
```
plt.figure(figsize=(15, 10))

sns.set(font_scale=1.5)
sns.set_style('ticks') # change background to white background

ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+', scatter_kws={'s': 200})
ax.set(xlabel='Year', ylabel='Total Immigration')
ax.set_title('Total Immigration to Canada from 1980 - 2013')
```

Out[51]:

```
Text(0.5, 1.0, 'Total Immigration to Canada from 1980 - 2013')
```



Or to a white background with gridlines.

In [52]:

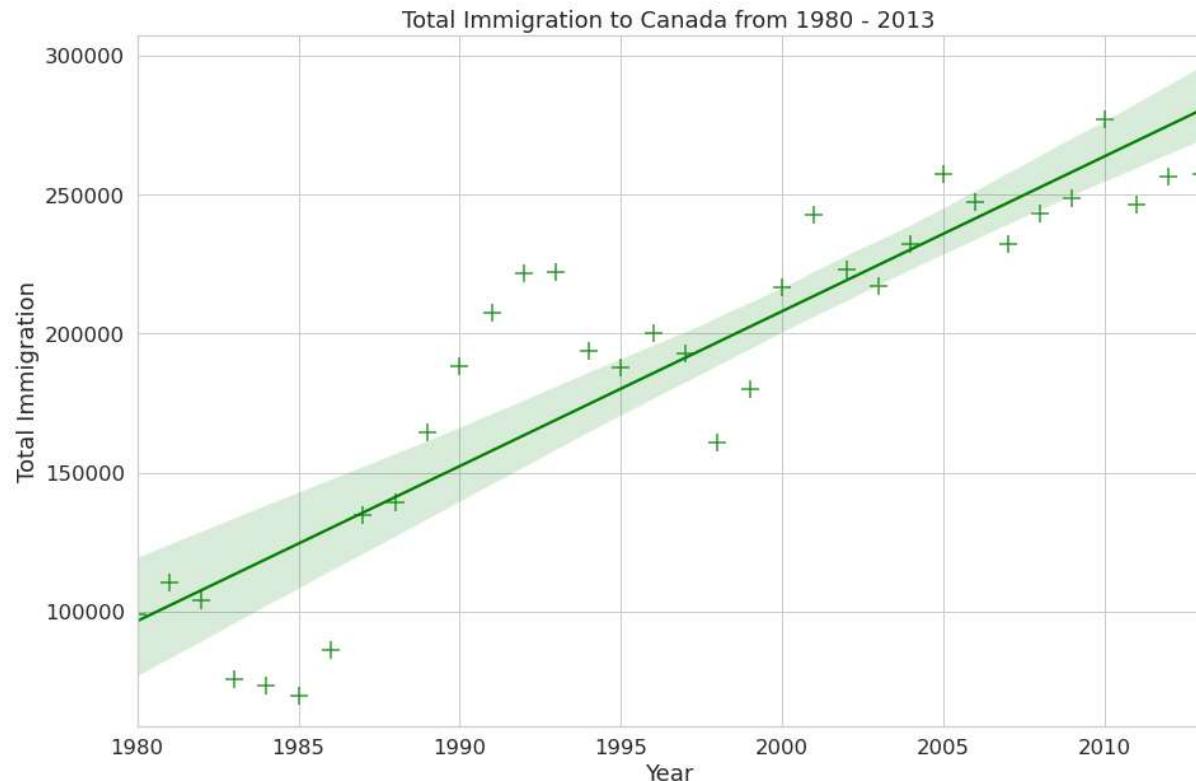
```
plt.figure(figsize=(15, 10))

sns.set(font_scale=1.5)
sns.set_style('whitegrid')

ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+', scatter_kws={'s': 200})
ax.set(xlabel='Year', ylabel='Total Immigration')
ax.set_title('Total Immigration to Canada from 1980 - 2013')
```

Out[52]:

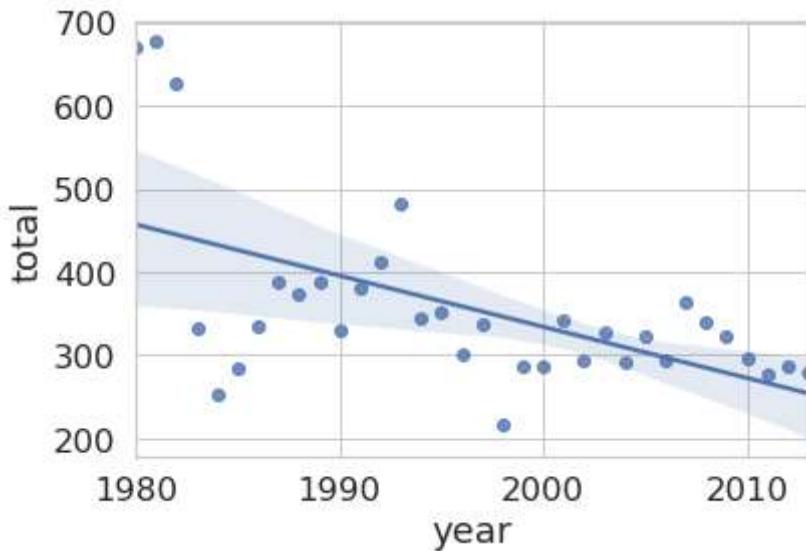
Text(0.5, 1.0, 'Total Immigration to Canada from 1980 - 2013')



**Question:** Use seaborn to create a scatter plot with a regression line to visualize the total immigration from Denmark, Sweden, and Norway to Canada from 1980 to 2013.

In [60]:

```
### type your answer here
df_countries = df_can.loc[['Denmark', 'Norway', 'Sweden'], years].transpose()
df_countries.head()
df_total = pd.DataFrame(df_countries.sum(axis=1))
df_total.reset_index(inplace=True)
df_total.columns = ['year', 'total']
df_total
df_total['year'] = df_total['year'].astype(int)
ax = sns.regplot(x='year', y='total', data=df_total)
```



► Click here for a sample python solution

**Thank you for completing this lab!**

## Author

[Alex Aklson \(<https://www.linkedin.com/in/aklson/>\)](https://www.linkedin.com/in/aklson/)

## Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2021-01-21	2.2	Lakshmi Holla	Updated TOC markdown cell
2020-11-03	2.1	Lakshmi Holla	Changed URL of excel file
2020-08-27	2.0	Lavanya	Moved lab to course repo in GitLab

**© IBM Corporation 2020. All rights reserved.**